

Artificial Intelligence

Learning Chapter 18

2007-10-29

Erik Billing
Umeå University, Computing Science

Learning

Learning is essential for unknown environments,
i.e., when designer lacks omniscience

Learning is useful as a system construction method,
i.e., expose the agent to reality rather than trying to write
it down

Learning modifies the agent's decision mechanisms to
improve performance

2007-10-29

Erik Billing
Umeå University, Computing Science

Learning

The agents percepts should be used not only for acting, but
also for improving future performance.

Tasks to learn for an agent:

- What state will be the result of an action?
- How will the changing world evolve?
- What is the value of each state
- Which kind of states has high (low) value?
- Which percepts are relevant?

Learning task – estimations of functions $y=f(x)$: $x \rightarrow y$

2007-10-29

Erik Billing
Umeå University, Computing Science

Types of Learning

Supervised learning:

Given a value x , $f(x)$ is immediately provided by a
"supervisor", $f(x)$ is learned from a number of
examples: $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$

Reinforcement learning:

A correct answer y is not provided for each x .
Rather a general evaluation is proved after a sequence of actions
(occasional rewards)

Unsupervised learning:

The agent learns relationships among its percepts. i.e. it perform
clustering.

2007-10-29

Erik Billing
Umeå University, Computing Science

Inductive Learning

The agent is fed with examples $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$

Problem: find a hypothesis h
such that $h \approx f$
given a training set of examples

Which hypothesis is best?

All learning methods make assumptions about f .

This preference is called a "bias".

h should perform well on the examples AND on unseen
examples: "h should generalise well".


2007-10-29

Erik Billing
Umeå University, Computing Science

Inductive Learning Method

Construct/adjust h to agree with f on training set
(h is consistent if it agrees with f on all examples)

E.g., curve fitting:



2007-10-29


Erik Billing
Umeå University, Computing Science



Inductive Learning Method

Construct/adjust h to agree with f on training set
(h is consistent if it agrees with f on all examples)

E.g., curve fitting:



2007-10-29


Erik Billing
Umeå University, Computing Science



Inductive Learning Method

Construct/adjust h to agree with f on training set
(h is consistent if it agrees with f on all examples)

E.g., curve fitting:



2007-10-29


Erik Billing
Umeå University, Computing Science



Inductive Learning Method

Construct/adjust h to agree with f on training set
(h is consistent if it agrees with f on all examples)

E.g., curve fitting:



2007-10-29


Erik Billing
Umeå University, Computing Science



Inductive Learning Method

Construct/adjust h to agree with f on training set
(h is consistent if it agrees with f on all examples)

E.g., curve fitting:



2007-10-29


Erik Billing
Umeå University, Computing Science



Inductive Learning Method

Construct/adjust h to agree with f on training set
(h is consistent if it agrees with f on all examples)

E.g., curve fitting:



Occham's razor: maximize
a combination of consist-
ency and simplicity

2007-10-29

Erik Billing
Umeå University, Computing Science



Attribute-Based Representations

Examples described by attribute values (Boolean,
discrete, continuous, etc.)

E.g., situations where I will/won't wait for a table:


Example	Attributes										Target
	Alt	Bar	Fri	Hum	Pat	Price	Rain	Res	Type	Est	
X_1	T	F	F	T	Some	\$\$\$	F	F	French	0-10	T
X_2	T	F	F	T	Some	\$	F	F	Thai	30-60	F
X_3	F	T	F	F	Some	\$	F	F	Burger	0-10	T
X_4	T	F	T	T	Full	\$	F	F	Thai	10-30	T
X_5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X_6	F	T	F	T	Some	\$	T	T	Italian	0-10	T
X_7	F	T	F	F	None	\$	T	F	Burger	0-10	F
X_8	F	F	F	T	Some	\$	T	T	Thai	0-10	T
X_9	F	T	T	F	Full	\$	T	F	Burger	>60	F
X_{10}	T	T	T	T	Full	\$\$\$	F	T	Italian	10-30	F
X_{11}	F	F	F	F	None	\$	F	F	Thai	0-10	F
X_{12}	T	T	T	T	Full	\$	F	F	Burger	30-60	T

2007-10-29

Erik Billing
Umeå University, Computing Science

Decision tree

One possible representation for hypotheses
E.g., here is the "true" tree for deciding whether to wait:




2007-10-29

Erik Billing
Umeå University, Computing Science

Expressiveness

Decision trees can express any function of the input attributes.
E.g., for Boolean functions, truth table row → path to leaf:



Trivially, \exists a consistent decision tree for any training set with one path to leaf for each example (unless f non-deterministic in x) but it probably won't generalize to new examples

Prefer to find more compact decision trees

2007-10-29

Erik Billing
Umeå University, Computing Science

Learning Decision Trees

The tree can often represent a set of examples in a compact way by identifying patterns in the examples.

This is a "simpler" function with hopefully better generalisation.


Decision trees is a major tool for Machine Learning in real applications.

2007-10-29

Erik Billing
Umeå University, Computing Science

Choosing an attribute

Idea: a good attribute splits the examples into subsets that are (ideally) "all positive" or "all negative"



Patrons? is a better choice – gives information about the classification

2007-10-29

Erik Billing
Umeå University, Computing Science

Information

Information answers questions

The less information we have about the answer initially, the more information is contained in the answer we get.

Information content I can be computed based on the prior probabilities P_i to P_n :

$$I(P_1, \dots, P_n) = -\sum_{i=1}^n P(v_i) \log_2 P(v_i)$$

The information content is also referred to as *entropy*.

2007-10-29

Erik Billing
Umeå University, Computing Science

Information

- Before asking any question, we can only base our answer on the prior probabilities.
- Since there is 6 positive and 6 negative examples in our training data, we need exactly one bit of information to answer the question if we should stay or not.

$$I_{\text{Wait}}\left(\frac{p}{p+n}, \frac{n}{p+n}\right) = \left[-\frac{6}{12} \log_2 \frac{6}{12}\right] + \left[-\frac{6}{12} \log_2 \frac{6}{12}\right] = [-0.5 * -1] + [-0.5 * -1] = 1$$

2007-10-29

Erik Billing
Umeå University, Computing Science



Information

- Suppose that we test the *Patrons* attribute.
- The information still required after testing *Patrons* is the average over the three alternatives *None*, *Some* and *Full*.

$$I_{None}\left(\frac{0}{2}, \frac{2}{2}\right) = \left[-\frac{0}{2}\log_2 \frac{0}{2}\right] + \left[-\frac{2}{2}\log_2 \frac{2}{2}\right] = [-0 * Inf] + [-1 * 0] = 0$$

$$I_{Some}\left(\frac{4}{4}, \frac{0}{4}\right) = \left[-\frac{4}{4}\log_2 \frac{4}{4}\right] + \left[-\frac{0}{4}\log_2 \frac{0}{4}\right] = [-1 * 0] + [-0 * Inf] = 0$$

$$I_{Full}\left(\frac{2}{6}, \frac{4}{6}\right) = \left[-\frac{2}{6}\log_2 \frac{2}{6}\right] + \left[-\frac{4}{6}\log_2 \frac{4}{6}\right] = [-0.33 * -1.59] + [-0.66 * -0.59] = 0.92$$

2007-10-29

Erik Billing
Umeå University, Computing Science



Information Theory

The information still required after testing attribute *A* is called the *reminder R*:

$$R(A) = \sum_{i=1}^r \frac{p_i + n_i}{p + n} * I\left(\frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i}\right)$$

Now, the *information gain G* from testing attribute *A* can be computed as:

$$G(A) = I\left(\frac{p}{p + n}, \frac{n}{p + n}\right) - R(A)$$

Choose the attribute *A* with the largest gain!

2007-10-29

Erik Billing
Umeå University, Computing Science



Performance of the Learning Algorithm

The critical question:

How does the hypothesis perform on unseen examples?

Test methodology:

- 1) Collect a large set of examples
- 2) Divide them randomly a training set and a test set
- 3) Generate a hypothesis using the training set
- 4) Measure the performance on examples from the test set. I.e. The percentage of examples in the test set that are correctly classified.
- 5) Optionally: repeat 1-4 many times

2007-10-29


Erik Billing
Umeå University, Computing Science



Performance Measurement

How do we know that $h \approx f$?

Try *h* on a new test set of examples
(use same distribution over example space as training set)



Learning curve: % correct on test set as a function of training set size

2007-10-29

Erik Billing
Umeå University, Computing Science



Noise and Overfitting

Often there are **TOO MANY** attributes !

It is hard to know in advance which attributes are relevant for the classification task. Irrelevant attributes act as noise. The problem is an example of OVERFITTING.

Makes the tree fit the training data
Gives BAD generalisation

We need help to determine when to stop adding nodes !
(all learning techniques have the general problems with overfitting)

2007-10-29

Erik Billing
Umeå University, Computing Science



Summary

- Learning needed for unknown environments.
- Learning method depends on available feedback, type of component to be improved, and its representation.
- For supervised learning, the aim is to find a simple hypothesis approximately consistent with training examples.
- Decision tree learning using information gain.
- Learning performance = prediction accuracy measured on test set.

2007-10-29

Erik Billing
Umeå University, Computing Science

Neural Networks

Two major types:

- Neural Networks**
In our brains
- Artificial Neural Networks**
Attempts to imitate our Biological Neural Networks with software and hardware

2007-10-29 Erik Billing
Umeå University, Computing Science

Why Would We Want to Imitate the Brain Architecture?

- Superior performance**
- Fault tolerant!**
Brain cells die all the time!
Still works because of local connections
- Can learn by examples (inductive learning)**
- Good generalization**
We manage to do something even with previously unseen information or in totally new situations!
- Can we imitate the brain ?**

2007-10-29 Erik Billing
Umeå University, Computing Science

Computing Elements

Neurons

- The computational unit of the brain
- Connected through Synapses which can be Excitatory or Inhibitory
- Computes an output as a function of all inputs from surrounding neurons

Units (nodes)

- The computational unit of an Artificial Neural Network
- Connected through links with numeric weights
- Computes an output as a function of all inputs from surrounding nodes

2007-10-29 Erik Billing
Umeå University, Computing Science

Neuron

Terminology:

- g : activation function
- a_i : activation level of node i (output of the node).

The **neuron** is also called **node**, or **unit**.

2007-10-29 Erik Billing
Umeå University, Computing Science

Activation Functions

The output from a node:

$$a_i = g\left(\sum_{j=1}^n W_{j,i} a_j\right)$$

Common activation functions g :

- Step function:** $\text{step}_t(x) = \begin{cases} 1, & x \geq t \\ 0, & x < t \end{cases}$
- Sign function:** $\text{sign}(x) = \begin{cases} +1, & x \geq 0 \\ -1, & x < 0 \end{cases}$
- Sigmoid:** $\text{sigmoid}(x) = \frac{1}{1+e^{-x}}$
- tanh:** $\tanh(x)$

2007-10-29 Erik Billing
Umeå University, Computing Science

Perceptron

Input layer with 4 inputs
Output layer with 3 output nodes

2007-10-29 Erik Billing
Umeå University, Computing Science

What Functions Can a Perceptron Represent?

$a = \text{step}\left(\sum_{j=0}^n W_j a_j\right)$

a_1
 $W_1 = 1$
 a_2
 $W_2 = 1$

$W_0 = 1.5$

$a = \text{step}(-W_0 + W_1 a_1 + W_2 a_2) = \text{step}(-1.5 + a_1 + a_2)$
AND

a_1
 $W_1 = 1$
 a_2
 $W_2 = 1$

$W_0 = 0.5$

$a = \text{step}(-W_0 + W_1 a_1 + W_2 a_2) = \text{step}(-0.5 + a_1 + a_2)$
OR

a_1
 $W_1 = -1$

$W_0 = -0.5$

$a = \text{step}(-W_0 + W_1 a_1) = \text{step}(0.5 - a_1)$
NOT

2007-10-29 Erik Billing Umeå University, Computing Science

Linear Separability

Consider a perceptron with two inputs and a threshold (bias):

- The perceptron fires if $w_1 a_1 + w_2 a_2 - t \geq 0$
- Recall the weights for the “and” perceptron:
 - $a_1 + a_2 - 1.5 \geq 0$
- This is really the equation for a line!
 - $a_2 \geq -a_1 + 1.5$

The activation threshold for a perceptron is actually a linearly separable “hyperplane” in the space of inputs

2007-10-29 Erik Billing Umeå University, Computing Science

Linear Separability

$I_1 \wedge I_2$

$I_1 \vee I_2$

$\neg I_1$

$I_1 \oplus I_2$

2007-10-29 Erik Billing Umeå University, Computing Science

Multi-Layer Networks

The structure of a multi-layer network is fairly straightforward:

- The **input layer** is the set of features (percepts)
- Next is a **hidden layer**, which has an arbitrary number of perceptrons called hidden units that take the features (input layer) as inputs
- The perceptron(s) in the **output layer** then takes the outputs of the hidden units as its inputs

2007-10-29 Erik Billing Umeå University, Computing Science

Multi-Layer Network

2007-10-29 Erik Billing Umeå University, Computing Science

Training Multi-Layer Network

Training multi-layer networks can be a bit complicated (the weight space is large!)

- The perceptron rule works fine for a single unit that mapped input features to the final output value
- But hidden units do not produce the final output
- Output unit(s) take other perceptrons – not known feature values – as inputs

The solution is to use the **back-propagation algorithm**, which is an intuitive **extension of the perceptron training algorithm**

2007-10-29 Erik Billing Umeå University, Computing Science



Perceptron Training

Conceptually, the perceptron rule does:

- Compare the perceptron's output (s) to what it *should* have been (t, e.g. *true*), i.e. compute error
- If the error is large, assign "blame" to the weight/input combinations that most influenced the wrong call, and raise/lower the weights accordingly
- If the error is small, don't change them as much

The key parameter is the learning rate η :

- If too small, learn slowly and convergence takes forever
- If too large, can make changes that are too drastic

2007-10-29

Erik Billing
Umeå University, Computing Science



Perceptron Learning

- A perceptron learns by adjusting its weights in order to minimize the error on the training set
- Consider updating the value for a single weight on a single example with the perceptron learning rule:

$$Err = y - a_i$$

$$x_j = a_j \times w_j$$

$$\nabla w_j = Err \times x_j \times \alpha$$

$$w_j \leftarrow w_j + \nabla w_j$$

y : The desired output

a_i : Activation of current node i

Err : The output error

w_j : Weight between node j and i

x_j : Weight influence from input j

2007-10-29

Erik Billing
Umeå University, Computing Science



Perceptron Learning

Now that we can update a single weight on one example, we want to update all weights so that we minimize error on the whole training set

- So this is really an optimization search in weight space, which is the hypothesis space for perceptrons

For reasons we won't get into, it's actually useful to minimize the squared error over the whole set:

- $E[w] = \frac{1}{2} \sum_d (\text{true}_d - a_d)^2$
- Where $E[w]$ is the sum of squared errors for the weight vector w , and d ranges over examples in the training set

2007-10-29

Erik Billing
Umeå University, Computing Science




Gradient Descent

Recall that minimization problems are called "gradient descent" tasks

If we have a perceptron with 2 weights, we want to find the pair of weights (i.e. point in 2D weight space) where $E[w]$ is the lowest

But the weights are continuous values, so how do we know how much to change them?



2007-10-29

Erik Billing
Umeå University, Computing Science



Gradient Descent

Solution: use calculus

- Compute the magnitude and direction of the gradient of the error surface by using a partial derivative:
 - $\nabla E[w] = [\delta E/\delta w_0, \delta E/\delta w_1, \dots, \delta E/\delta w_n]$
- We want to update each weight w_i by ∇w_i :
 - $\nabla w_i = -\eta [\delta E/\delta w_i]$
- And if we combine this with our weight update rule, we get the following complete perceptron training rule:
 - $w_i \leftarrow w_i + (-\eta [\delta E/\delta w_i])$
 - This makes sense: if ($\text{true} - o$) is positive, the weight should be increased for positive inputs a_i , and decreased for negatives

2007-10-29

Erik Billing
Umeå University, Computing Science



Back-Propagation (BP)

BP generalizes the perceptron rule:

- Gradient-descent search to minimize error on the training data (again, usually in iterative mode)
- In the forward pass, features are fed forward to the output layer where error is calculated

From earlier slide:

- The perceptron rule worked fine for a single unit that mapped input features to the final output value
- But hidden units do not produce the final output
- Output unit(s) take other perceptrons – not known feature values – as inputs

2007-10-29

Erik Billing
Umeå University, Computing Science



Problems with BP

Because BP is a gradient descent (hill-climbing) search, it suffers from the same problems:

- Doesn't necessarily find the globally best weight vector
 - Convergence is determined by the starting point (randomly initialized weights)
 - If η set too large, can "bounce" right over the global minimum into a local minimum

To deal with these problems:


- Usually initialize weights close to 0
- Can repeat training with multiple random restarts

2007-10-29

Erik Billing
Umeå University, Computing Science



Training and Over Training



Over training occurs when the model has fit all characteristics and starts fitting noise in data.

The problem increases with the number of weights (i.e. the model complexity)

2007-10-29

Erik Billing
Umeå University, Computing Science



Overfitting in Neural Nets

As usual, we can use a tuning set to avoid overfitting in neural nets

- We can train several candidate structures and use the tuning set to find one that's appropriately expressive

More common: given a network structure, use early stopping by evaluating the network on the tuning set after each epoch

- Stop when performance begins to dip on the tuning set
- Sometimes allow a fixed number of epochs beyond the dip... just in case it goes back up

2007-10-29

Erik Billing
Umeå University, Computing Science



Neural Network Characteristics

- Universal approximators – very powerful and can approximate any function, but they do have drawbacks
 - Many weights to learn: training can take a while
 - Sensitive to structure, initial weights, and learning rate
- Well suited for parallel computing since each node only depend on the connected neighbor nodes.
- Seem to have good generalization abilities. Tolerant to noise in the input.

2007-10-29

Erik Billing
Umeå University, Computing Science



Neural Network Characteristics

Can be applied to a number of totally different applications

"Black boxes". Hard to describe WHY the network produces a certain output or decision.

Prior knowledge. Hard to incorporate knowledge about the function into the architecture.

2007-10-29

Erik Billing
Umeå University, Computing Science



Summary

"The second best solution to pretty much everything!"

- Perceptrons are mathematical models of neurons (brain cells)
 - Learn linearly separable functions
 - Insufficiently expressive for many problems
- Neural Networks are machine learning models that have multiple layers of perceptrons
 - Trained using back-propagation, a gradient descent search through weight space (NN hypothesis space)
 - Sufficiently expressive for any classification or regression task, also quite robust to noise.

2007-10-29

Erik Billing
Umeå University, Computing Science



Summary

Many applications:

- Speech processing, driving, face/handwriting recognition, backgammon, checkers, etc.

Disadvantages:

- Overly expressive: prone to overfitting
- Difficult to design appropriate structure
- Many parameters to estimate: slow training
- Hill-climbing can get stuck in local optima
- Poor comprehensibility (black box-problem)

2007-10-29

Erik Billing
Umeå University, Computing Science